

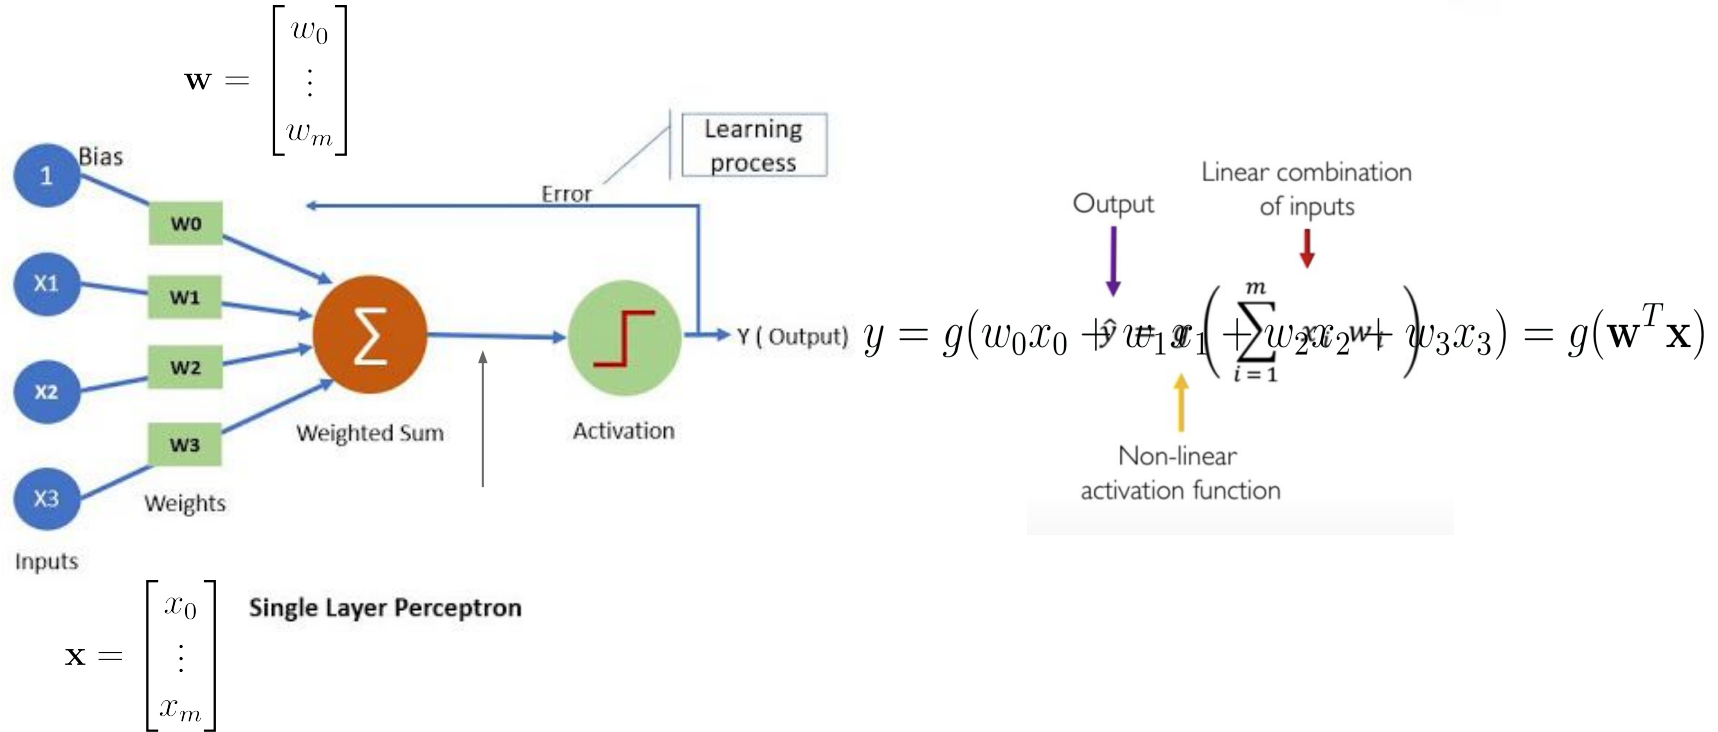


Deep Learning


1.1

Introduction To Deep Learning


Perceptron: Forward Propagation




Perceptron: Multi Output



Single Layer Neural Network




Multi-Layer Perceptron MLP



```
model = keras.Sequential([
    layers.Dense(4, input_shape=(3, ), activation="relu"),
    layers.Dense(4, activation="relu"),
    layers.Dense(1, activation="sigmoid"),
])
model.compile(loss="binary_crossentropy", optimizer="adam")
model.summary()
```

Multi-Layer Perceptron MLP



```
inputs = tf.keras.Input(shape=(3, ))  
l_1 = layers.Dense(4, activation='relu')(inputs)  
l_2 = layers.Dense(4, activation='relu')(l_1)  
outputs = layers.Dense(1, activation='sigmoid')(l_2)  
model = keras.Model(inputs=inputs, outputs=outputs)
```

Cost Functions

- Regression:

- **Mean squared error (MSE):**
$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$
- **Mean absolute error (MAE)**
$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |\hat{y}_i - y_i|$$

```
model.compile(loss='mse')
```

```
model.compile(loss='mae')
```

- Classification:

- **Binary Cross-Entropy (log-loss):**
$$\mathcal{L}(\hat{y}, y) = -\frac{1}{N} \sum_{i=1}^N y_i \cdot \log(\hat{y}_i) + (1 - y_i) \cdot \log(1 - \hat{y}_i)$$

```
model.compile(loss='binary_crossentropy')
```


- **Categorical Cross-Entropy:**
$$\mathcal{L}(\hat{y}, y) = -\frac{1}{N} \sum_{i=1}^N \sum_{k=1}^K y_{ik} \cdot \log(\hat{y}_{ik})$$

```
model.compile( loss='categorical_crossentropy')
```

```
model.compile( loss='sparse_categorical_crossentropy')
```




Regularization: Dropout



- During training, randomly set some activations to 0 with probability p .
- Not in prediction.

```
from tensorflow.keras.layers import  
Dropout  
model = Sequential()  
model.add(Dense(60, activation='relu',  
input_shape=input_shape))  
model.add(Dropout(0.2))  
model.add(Dense(30, activation='relu'))  
model.add(Dense(1, activation='sigmoid'))
```

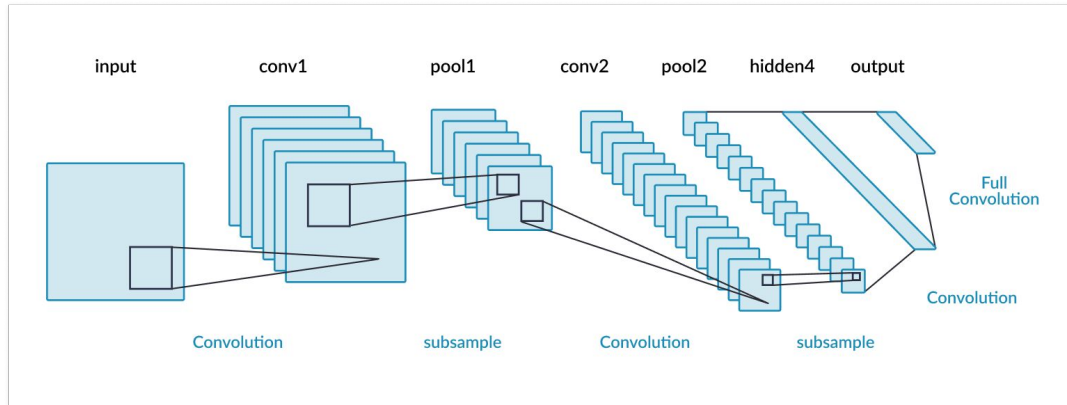
Regularization: Layer normalization




- Normalizes the activations along the feature direction instead of mini-batch direction
- BN normalizes each feature independently across the mini-batch. LN normalizes each of the inputs in the batch independently across all features.

```
from keras.layers.normalization import  
BatchNormalization  
model = Sequential()  
model.add(Dense(64, input_shape))  
model.add(LayerNormalization(axis=1))  
model.add(Activation('tanh'))  
model.add(Dropout(0.5))
```

CNN



CNN



CNN

```
image_size = (128,128,3)
inputs = tf.keras.Input(shape=image_size)

# Conv Layer 1
conv_1 = layers.Conv2D(4, 3, padding='valid', activation='relu')(inputs)
pool_1 = layers.MaxPooling2D(pool_size=(2, 2), name='pool_1')(conv_1)

# Conv Layer 2
conv_2 = layers.Conv2D(4, 3, padding='same', activation='relu')(pool_1)
pool_2 = layers.MaxPooling2D(pool_size=(2, 2))(conv_2)

# Flattening
flat = layers.Flatten(name='flatten')(pool_2)

# Fully-connected
dense = layers.Dense(64, activation='relu')(flat)
outputs = layers.Dense(5, activation='softmax')(dense)

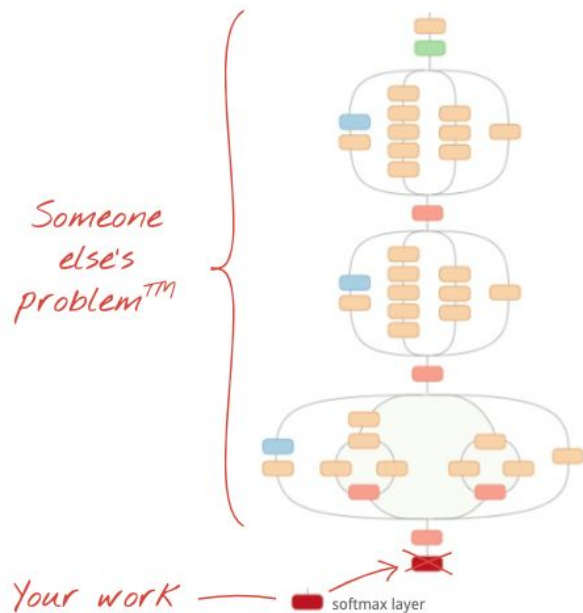
model = keras.Model(inputs=inputs, outputs=outputs)
```



CNN: Transfer Learning

```
from tensorflow.keras.applications import MobileNetV2
pretrained_model = MobileNetV2(input_shape=image_size,
                                include_top=False)
pretrained_model.trainable = False

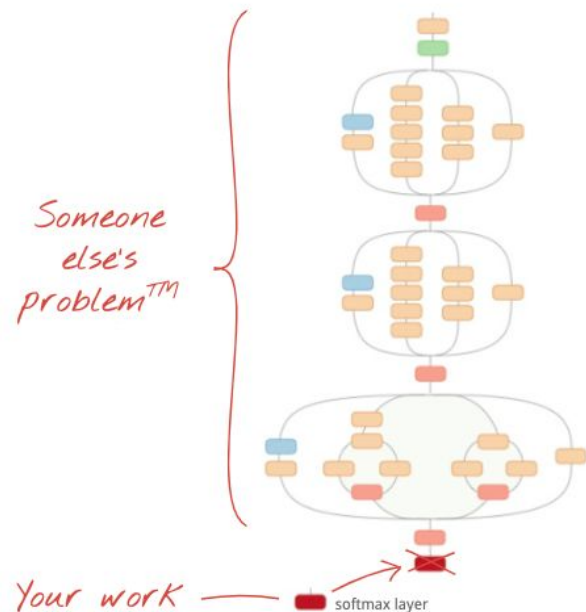
model_t1 = tf.keras.Sequential([
    pretrained_model,
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(5, activation='softmax')
])
```



CNN: Fine tuning

```
from tensorflow.keras.applications import MobileNetV2
pretrained_model = MobileNetV2(input_shape=image_size,
                                include_top=False)
pretrained_model.trainable = True

model_t1 = tf.keras.Sequential([
    pretrained_model,
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(5, activation='softmax')
])
```





UNIVERSIDAD
COMPLUTENSE
DE MADRID

